PROJECT ATTEMPT 2

Project Results

Review your submission outcome, score components, and peer feedback.

See all projects | Project statistics

TOTAL SCORE	STATUS	REVIEWED ENOUGH PEERS
28	Passed	Yes

Criteria Breakdown

Retrieval evaluation — 2

- No evaluation of retrieval is provided (0 points) — 0 votes
- Only one retrieval approach is evaluated (1 point) — 0 votes
- Multiple retrieval approaches are evaluated, and the best one is used (2 points) — 2 votes

RAG evaluation — 2

- No evaluation of RAG is provided (0 points) — 0 votes
- Only one RAG approach (e.g., one prompt) is evaluated (1 point) — 1 vote
- Multiple RAG approaches are evaluated, and the best one is used (2 points) — 1 vote

Interface — 2

- No way to interact with the application at all (0 points) — 0 votes
- Command line interface, a script, or a Jupyter notebook (1 point) — 0 votes
- UI (e.g., Streamlit), web application (e.g., Django), or an API (e.g., built with FastAPI) (2 points) — 2 votes

Ingestion pipeline — 2

- No ingestion (0 points) — 0 votes
- Semi-automated ingestion of the dataset into the knowledge base, e.g., with a Jupyter notebook (1 point) — 1 vote
- Automated ingestion with a Python script or a special tool (e.g., Mage, dlt, Airflow, Prefect) (2 points) — 1 vote

Monitoring — 2

- No monitoring (0 points) — 0 votes
- User feedback is collected OR there's a monitoring dashboard (1 point) — 0 votes
- User feedback is collected and there's a dashboard with at least 5 charts (2 points) — 2 votes

Containerization — 2

- No containerization (0 points) — 0 votes
- Dockerfile is provided for the main application OR there's a docker-compose for the dependencies only (1 point) — 0 votes
- Everything is in docker-compose (2 points) — 2 votes

Problem description — 2

- The problem is not described (0 points) — 0 votes
- The problem is described but briefly or unclearly (1 point) — 0 votes
- The problem is well-described and it's clear what problem the project solves (2 points) — 2 votes

RAG flow — 2

- No knowledge base or LLM is used (0 points) — 0 votes
- No knowledge base is used, and the LLM is queried directly (1 point) — 0 votes
- Both a knowledge base and an LLM are used in the RAG flow (2 points) — 2 votes

Reproducibility — 2

- No instructions on how to run the code, the data is missing, or it's unclear how to access it (0 points) — 0 votes
- Some instructions are provided but are incomplete, OR instructions are clear and complete, the code works, but the data is missing (1 point) — 0 votes
- Instructions are clear, the dataset is accessible, it's easy to run the code, and it works. The versions for all dependencies are specified. (2 points) — 2 votes

Best practices — 1

- Hybrid search: combining both text and vector search (at least evaluating it) (1 point) — 2 votes
- Document re-ranking (1 point) — 0 votes
- User query rewriting (1 point) — 0 votes

Bonus points — 0

- Deployment to the cloud (2 points) — 0 votes
- Bonus point (1 point) — 0 votes
- Bonus point (1 point) — 0 votes
- Bonus point (1 point) — 0 votes

Score Components

Project score	19
Peer review score	9
Project LiP	0
Peer review LiP	0
FAQ contribution	0

Submission

Project link
https://github.com/soumen7saha/comprehensive-medical-qna-assistant

Commit ID
de7f5a3

Submitted at
Aug. 15, 2026, 1:55 p.m.

Evaluation Feedback

PEER FEEDBACK

Excellent submission. Retrieval evaluation is thorough — comparing text, vector, and hybrid (RRF) search with hit rate/MRR on a 867-query ground truth set, and picking hybrid as the winner, is exactly the kind of rigor this criterion rewards. RAG evaluation judges good/bad but doesn't compare alternative prompts. Best practices: hybrid search only — re-ranking or query rewriting would be a natural next step. Nice that the dataset is bundled in the repo, which made reproducibility straightforward to confirm.

PEER FEEDBACK

Strong and well-structured RAG project. The README is clear and well written, with helpful architecture details, screenshots, and demo videos. Retrieval evaluation is particularly strong, comparing text, vector, and hybrid search and using the best-performing approach. The Streamlit UI, monitoring dashboard, feedback collection, and Docker setup are also well implemented. The main improvement would be evaluating multiple RAG approaches, such as different prompts or configurations, instead of only the final pipeline. The ingestion script could also be made directly executable as a complete ingestion pipeline.